

Computer Networks

First Laboratory Work

Gabriel Carvalho (up202208939@up.pt)

Vasco Melo (up202207564@up.pt)

Summary:

This project, carried out as part of the Computer Networks course unit, focuses on implementing a download program that follows the file transfer protocol, as well as establishing, configuring and utilizing a computer network.

Through this project, we were able to apply the knowledge gained in the theoretical lessons and consolidate it through the implementation of the program with the specified protocol and the network configuration.

Introduction:

The goal of this project was to develop and test a download program using FTP, as well as to configure a computer network, in accordance with the specifications outlined in the guidelines, with the final objective of transferring a file from the internet using the configured network.

This report is structured into three sections, each focusing on a specific aspect of the project:

1. **Downloader:** This section describes the architecture of the download application and the results we obtained from its implementation
2. **Network Configuration and Analysis:** This section includes the documentation, and log analysis for all the experiments made throughout the project
3. **Conclusions:** The final section synthesizes the information presented in the previous one, some reflections and conclusions as well as potential improvements.

1. Downloader:

Application Architecture

The objective of the first part of the project was to develop an application capable of performing file downloads using the FTP protocol. To accomplish this, the FTP protocol specification and pre-existing FTP applications were used to gain an understanding of socket programming and its implementation. All references consulted are documented in the reference section.

When the application is initialized, a URL is provided by the user. This URL follows a fixed structure, which makes the parsing process universal for any URL. Our URL parsing

begins by validating the URL format, then proceeds to identify the location of the URL separators (e.g., the "@" symbol separates the username and authentication password). The relevant fields required for file transfer are extracted in the following order: user and password for server authentication, which are set to default values (user: ftp, password: eu@) if not specified in the input; hostname, the server name where the communication will be established; port, the FTP server port, which defaults to 21 if not specified in the URL; filepath, the path to the desired file; and file, the name and extension of the file to be transferred. Finally, the server's IP address is obtained using the provided hostname and port, enabling the establishment of a socket connection to the server by the usage of `getaddrinfo` of the library `netdb.h`.

The application is based on socket programming, where it uses a stream socket as the communication medium between the application and the FTP server. The application sends commands (e.g., "USER") and the server responds with three-digit codes (e.g., 230 for successful login). To send commands to the FTP server, the `send` function is used, and to read the response, the `recv` function is used, both from the `sys/socket.h` library.

The communication between the application and the FTP server begins with the handshake process, during which the user's credentials are transmitted. This is accomplished using the "USER <username>" command to send the username and the "PASS <password>" command to transmit the password.

In the next phase, information is sent regarding how the download will be performed, including the file and its size. Authentication is then carried out, followed by sending the "PASV" command to put the server into Passive Mode. In response, the server sends the IP address and port number where it is listening, allowing the client to establish a connection. To do this, the client will create a second socket using the IP and port provided by the FTP server. Then, the "TYPE I" command is sent to inform the server that the file to be transferred is in binary format. Next, the "SIZE <filename>" command is sent to check the size of the desired file. Finally, the "RETR <file path>/<file name>" command is sent to the server to retrieve the <file name> file.

During the download phase, the second socket is used, created with the IP address and port provided by the FTP server after sending the "PASV" command. The FTP server will transmit the requested file content to the second socket, where the application will be listening to read the data in order to perform the download. The `write` function from the `sys/socket.h` library is used for this purpose. Once all bytes have been sent, the application reads the response sent by the FTP server through the first socket, which provides the download status (226 if the operation was successful).

Finally, the application sends the "QUIT" command to terminate the connection between the application and the FTP server. Both sockets are then closed, marking the completion of the file download process.

Tests and Results

To test the FTP application, various file transfer tests were conducted with different sizes and types of files. Additionally, connection attempts were made to both public servers

provided by Moodle and the FEUP FTP server, with the latter yielding a positive result. The tests covered both anonymous and authenticated modes. Finally, the FTP application was successfully used to carry out Experiment 6 of the second part of the project.

Below is an example of a successful transfer (the Wireshark capture from Experiment 6 of the second part of the project, and an image showing all commands and responses exchanged between the client and the server of the example shown below has been included in the appendices).

```
pchmelo@pchmelo-VirtualBox:~/Downloads/RCOM-project2-master$ ./download ftp://demo:password@test.rebex.net/readme.txt
Login successful
Transferring file...
Status: 100.00% [=====]
Transfer complete!
pchmelo@pchmelo-VirtualBox:~/Downloads/RCOM-project2-master$ cat readme.txt
Welcome to test.rebex.net!

You are connected to an FTP or SFTP server used for testing purposes
by Rebex FTP/SSL or Rebex SFTP sample code. Only read access is allowed.

For information about Rebex FTP/SSL, Rebex SFTP and other Rebex libraries
for .NET, please visit our website at https://www.rebex.net/

For feedback and support, contact support@rebex.net

Thanks!
```

Note: This image shows the download being made with the debug mode turned off. An image showing the process with the debug mode on is shown [here](#).

2. Network Configuration and Analysis:

Experiment 1 - Configure an IP Network

This experiment's goal was to configure two IP addresses of two different computers, tux13 and tux14, both connected to a switch. We were also tasked with analyzing how the ARP protocol worked, as well as its behaviour when configured IP addresses are deleted from an ARP table.

In order to achieve this, we started by connecting the E1 of each tux to the switch and used the command ifconfig to configure their respective IPs. After which, we used the ifconfig command again, but this time to visualize the IP and MAC addresses of each respective computer, which were respectively 172.16.10.1/24 and 00:50:fc:ed:bb:37 for tux13, and 172.16.10.254/24 and 00:c0:df:04:20:99 for tux14.

So as to verify connectivity between the two computers we used the ping command which generates and sends via broadcast ARP packets, a type of network protocol used to map an IP address to a MAC address in a local network, effectively connecting the data link layer(MAC) to the network layer(IP), until obtaining the destination MAC address, then sends ICMP Echo Request packets to initiate the communication. If the target device is reachable, it responds with ICMP Echo Reply packets. These ICMP packets allow the ping command to test connectivity between the two devices, measure round-trip time, and check if the target device responds correctly.

Another great way of troubleshooting, and testing network connectivity locally within a system is the loopback interface, which is a virtual network interface in a computer system

that allows a device to communicate with itself through its assigned IP address (usually 127.0.0.1 in IPv4).

During the ping command, when sending ARP packets, after a connection is established, it gets recorded the ARP table of both computers in order to make communication faster and more efficient by allowing devices to quickly resolve IP addresses to MAC addresses without requiring new ARP exchanges for every packet sent. When these entries are deleted, an exchange needs to be made again to restore the associations.

Additional information from this experiment, such as the size or type of a received ethernet frame can be seen in the Wireshark capture [here](#), and the commands used [here](#).

Experiment 2 - Implement two bridges in a switch

This experiment's goal was the creation of two local networks (LAN) through the use of bridges in the switch, one containing tux13 and tux14, and the other containing just tux12.

To do this, we began from where we stopped in experiment 1, connected the E1 of tux12 to the switch and used the ifconfig command to configure tux12 with the IP 172.16.11.1/24. In order to create the bridges, we must configure the switch by connecting its console with any tux's serial port. After which, we use the command '/interface bridge add' to create bridge0 and bridge1, use '/interface bridge port remove' to remove the ports each tux's interface was connected to by default, and finally used '/interface bridge port add' to add the tuxes 13 and 14 to bridge0 and tux12 to bridge1.

In this experiment, there were two broadcast domains, each corresponding to one of the two local networks we had implemented. When broadcasting from tux13, tux14 could be reached but tux12 couldn't. This happened because tux13 and 14 were on the same LAN, whereas tux12 resides in an isolated LAN. Doing the same in tux12 showed that both other tuxes were unreachable, due to the same reason stated previously.

The conclusions we reached in the last paragraph were based on the analyses of the Wireshark captures. We observed that in tux13's broadcast, both ICMP requests and replies were present. However, in tux12's broadcast, only ICMP requests were seen, with no replies.

Additional information from this experiment can be seen in the Wireshark captures [here](#), and the commands used [here](#).

Experiment 3 - Configure a Router in Linux

This experiment's goal was to transform tux14 into a router, so as to tux13 and tux12 will be able to communicate through tux14, despite being on different LANs.

In order to complete our goal, we began from where we last stopped on experiment 2 and connected tux4's E2 to the switch, making sure to connect it to bridge1 using '/interface bridge port add'. After which, we used ifconfig to make its IP address 172.16.11.253/24, activated IP forwarding, and disabled ICMP echo ignore broadcasts. Lastly, using the 'route add' command we created new routes in tux12 and 13, using tux14 as a gateway, accessible by them both through 172.16.11.253/24 and 172.16.10.254/24 respectively.

Now, in tux13, after utilizing the ping command to tux12, all the ARP and ICMP requests would return a corresponding reply. The ARP and ICMP packets processed by Tux54 contained the destination machine's IP address but the MAC address of Tux54 itself, since it, like a router, is responsible for redirecting information between the two networks it has created. The routing tables generated through the creation of routes ensure that for every destination IP, there is another IP address (gateway) to which the source machine should forward the information.

Additional information from this experiment can be seen in the Wireshark captures [here](#), and the commands used [here](#).

Experiment 4 - Configure a Commercial Router and Implement NAT

This experiment's goal was to configure a commercial router with implemented NAT, and add it to bridge1 with the objective of obtaining access to the internet.

To do this, we started from where we last stopped in experiment 3 and connected ether1 and ether2 of the commercial router to PY.12 and the switch respectively. Later, we added the interface to bridge1 in the same way as stated in previous experiments, and removed the cable connected to the switch's console, connecting it to routerMT, so as to configure the router. On the router's console, we configured the IPs corresponding to each interface with the command '/ip address add'. Lastly, we created default routes in each tux connecting them to the router, and another route that would allow the router access to both LANs through tux14.

After removing the route connecting Tux12 to Tux14 and disabling ICMP redirects, when pinging tux13 from 12, the data packets were forwarded through the router by the default route until reaching the destination IP. In a second test, after enabling redirects once again, the packets didn't pass through the router due to a more direct route, this being to simply use the tux14, connecting the two LANs. The tests allowed us to conclude that ICMP packets follow the path of least resistance, performing redirects only when necessary.

We were also asked in this experience to ping an FTP server with and without NAT enabled, using the command '/ip firewall nat enable/disable 0', and observe the results. NAT (Network Address Translation) is a method that modifies the source or destination IP while they are in transit across a routing device, such as a router or firewall. It is used to enable multiple devices on a private network to access the internet, by translating their private IP addresses, used only in the local network, into a public IP address that is accessible over the internet. This allows devices to address the issues stemming from the limited number of addresses allowed by the ipv4 protocol.

With NAT enabled we could verify that a connection to the internet was established. Whereas, when we disabled it, a connection couldn't be established, due to the router not being able to translate the destination address received.

Additional information from this experiment can be seen in the Wireshark captures [here](#), and the commands used [here](#).

Experiment 5 - DNS

The goal of this experiment was to configure the DNS (Domain Name System), so as to, using the network we configured throughout these experiments, be able to access the internet and ping a website.

In order to achieve this, all we needed to do was change the content of the file '/etc/resolv.conf', in each tux, to include the following line: 'nameserver 10.227.20.3'. Now we were able to reach the internet and successfully pinged google.

In the Wireshark capture logs we observed that the DNS packets were the first ones being exchanged and happened before any other protocol. The reason for this is that all the protocols down the line will need the domain's IP address in order to work, and to obtain that address, the DNS is needed.

Additional information from this experiment can be seen in the Wireshark captures [here](#), and the commands used [here](#).

Experiment 6 - TCP connections

The goal of this experiment is to combine both the downloader, made in the first part of the project, and the configured network, established in the second, in order to analyse the network traffic captured during FTP file transfer so as to evaluate packet transmission, internal packet structure, and TCP's flow control and congestion control mechanisms.

By examining the Wireshark capture, we observed multiple TCP connections established for FTP operations. The initial connection setup showed the TCP handshake with SYN-ACK packets establishing the control channel on port 21. The capture revealed that the FTP application establishes two distinct TCP connections: a control connection for sending commands and a separate data connection for file transfer.

The FTP control connection carries commands and responses between client and server, as evidenced by messages such as "ProFTPD Server (Debian)" responses and user authentication exchanges. These control messages operate over a dedicated TCP connection, maintaining separation from the data transfer channel.

All observed TCP connections implement ARQ (Automatic Repeat Request) mechanisms, which ensure reliable data delivery by requiring receivers to acknowledge received packets and senders to retransmit lost packets. The capture shows this through sequence numbers (Seq), acknowledgment numbers (Ack), and window size fields (Win). For example, we observed consistent acknowledgment patterns with sequence numbers incrementing properly and corresponding acknowledgments confirming received data. The TCP connections maintained window sizes around 64256 bytes, indicating stable flow control conditions.

The flow control mechanism, which prevents a fast sender from overwhelming a slow receiver by allowing the receiver to specify how much data it can handle, is visible through the Window field in TCP headers. Throughout the capture, window size values remained

relatively stable, indicating that receiver buffer capacity was not a limiting factor during this transfer.

Regarding congestion control, which prevents network congestion by regulating the rate at which packets are sent, the capture shows TCP connections operating with Maximum Segment Size (MSS) values of 1460 bytes, as seen in the initial SYN packets. The relatively stable window sizes and consistent packet transmission patterns suggest that the network remained in a steady state during the capture period, without entering congestion scenarios that would trigger congestion avoidance mechanisms.

The transfer patterns visible in the capture align with TCP's standard congestion control mechanisms, though we don't see dramatic congestion window adjustments in this particular sample. This suggests the network conditions were favorable during the capture period, maintaining stable throughput without triggering congestion responses.

When examining the interaction between multiple TCP connections, the capture shows concurrent control and data connections operating without significant interference. The separation of control and data channels allows for efficient management of the FTP session while maintaining optimal data transfer rates. The consistent window sizes and acknowledgment patterns across both connections indicate that the network adequately handled the concurrent TCP streams without notable performance degradation.

Additional information from this experiment can be seen in the Wireshark captures [here](#), and the commands used [here](#).

3. Conclusions:

Through this project, we gained comprehensive insight into the intricate workings of network protocols and data transfer mechanisms in computer networks. The practical experiments enhanced our understanding of how data packets propagate through the network layer, including routing decisions and addressing schemes, and their subsequent transition to the link layer for physical transmission. This hands-on experience with protocol analysis and network traffic observation helped bridge the gap between theoretical networking concepts and their real-world implementation, providing valuable practical knowledge of modern network communications.

The project particularly illuminated the complex interactions between different networking layers, demonstrating how higher-level protocols like FTP rely on the robust mechanisms like TCP and IP, and how these protocols work together to ensure reliable and efficient data transfer across network infrastructure.

Annexes

References:

- <https://datatracker.ietf.org/doc/html/rfc959>
- https://github.com/thevickypedia/ftp_server_client
- <https://github.com/literallysofia/feup-rcom/tree/master>
- https://beej.us/guide/bgnet/pdf/bgnet_a4_c_2.pdf

Download Application Code:

download.h

C/C++

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <signal.h>
#include <netdb.h>
#include <string.h>
#include <strings.h>
#include <sys/stat.h>
#include <fcntl.h>

#define DEBUG 0
#define RESP_BUFFER_SIZE 1024
#define FILE_BUFFER_SIZE 1024

typedef struct connectionInfo{
    char user[128];
    char password[128];
    char hostname[128];
    char port[8];
    char filepath[256];
    char filename[128];
}

connectionInfo;

int parseUrl(char * url, connectionInfo* connInfo);
void ftpClient(connectionInfo cInfo);
void dataConnection(int sd2, char * filename, int filesize);
```

```
void loadingBar(float file_size_processed, float file_total_size);
```

download.c

```
C/C++
#include "download.h"

/*
Protocolo FTP
1. Parse URL
2. Procura pelas resolved addresses com as informações do hostname e port
3. Cria uma socket com as características dos resolved addresses
4. Conecta a socket ao servidor
5. Recebe a resposta do servidor e espera pelo código 220
6. Envia o comando USER com o username e espera pelo código 331
7. Envia o comando PASS com a password e espera pelo código 230
8. Envia o comando PASV e espera pelo código 227
9. Com o IP e porta do PASV, cria uma nova socket e conecta ao servidor
10. Envia o comando TYPE I e espera pelo código 200
11. Envia o comando SIZE com o nome do ficheiro e espera pelo código 213, e
    guarda o tamanho do ficheiro
12. Envia o comando RETR com o nome do ficheiro e espera pelo código 150v ou
    125
13. Recebe o ficheiro pela segunda socket
14. Fecha a segunda socket
15. Recebe a resposta do servidor e espera pelo código 226, para confirmar a
    transferência
16. Fecha a primeira socket
*/



int main(int argc, char** argv)
{
    if (argc!=2)
    {
        printf("Wrong number of arguments\n");
        return -1;
    }

    connectionInfo cInfo;

    if( parseUrl(argv[1], &cInfo) != 0 )
        printf("Wrong url input.\n");

    if (DEBUG){
        printf("\n<PARSE URL>\n");
    }
```

```

        printf("user = %s\n", cInfo.user);
        printf("password = %s\n", cInfo.password);
        printf("hostname = %s\n", cInfo.hostname);
        printf("port = %s\n", cInfo.port);
        printf("filepath = %s\n", cInfo.filepath);
        printf("filename = %s\n\n", cInfo.filename);
    }

    ftpClient(cInfo);

    return 0;
}

int sendComand(int sd, char * userCmd, char * response)
{
    if(send(sd, userCmd, strlen(userCmd), 0) == -1){
        printf("Error sending USER cmd\n");
        exit(-1);
    }

    sleep(1);

    int received = recv(sd, response, RESP_BUFFER_SIZE, 0);
    if(received == -1){
        perror("Error receiving response from USER");
        exit(1);
    }

    response[received] = 0;

    return received;
}

void ftpClient(connectionInfo cInfo)
{
    //--- GET IP ---

    struct addrinfo hints;
    struct addrinfo *result;

    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;                                     //IPv4
    hints.ai_socktype = SOCK_STREAM;                                //TCP
    hints.ai_flags = 0;

    //no flags
    hints.ai_protocol = 0;
    //Any protocol can be used
}

```

```

//getaddrinfo returns the resolved addresses
int r = getaddrinfo(cInfo.hostname, cInfo.port, &hints, &result);

if (r != 0) {
    fprintf(stderr, "getaddrinfo failed: %s\n", gai_strerror(r));
    exit(-1);
}

//--- SOCKET ---

//cria uma socket para o servidor ftp onde podem ser enviados comandos
int sd = socket(result->ai_family, result->ai_socktype,
result->ai_protocol);
if(sd == -1){
    perror("Error opening socket");
    exit(-1);
}

//--- CONNECT ---

//make the socket connect to the server using the resolved address
if(connect(sd, result->ai_addr, result->ai_addrlen) == -1){
    perror("Error connecting to server");
    exit(-1);
}

char response[RESP_BUFFER_SIZE];

int received = recv(sd, response, RESP_BUFFER_SIZE, 0);
if(received == -1){
    perror("Error receiving response from connect");
    exit(1);
}
response[received] = 0;
if (DEBUG) printf("<CONNECT RESPONSE>\n%s\n", response);

if(strncmp("220", response, 3) != 0){
    printf("Connection failed\n");
    exit(-1);
}

freeaddrinfo(result);

//--- LOGIN ---

char userCmd[128];
char passCmd[128];

```

```

sprintf(userCmd, "USER %s\r\n", cInfo.user);
if (DEBUG) printf("<USER CMD>\n%s\n", userCmd);

received = sendComand(sd, &userCmd, &response);
if (DEBUG) printf("<USER RESPONSE>\n%s\n", response);

if(strncmp("331", response, 3) != 0){
    printf("Error in USER\n");
    exit(-1);
}

sprintf(passCmd, "PASS %s\r\n", cInfo.password);
if (DEBUG) printf("<PASS CMD>\n%s\n", passCmd);

received = sendComand(sd, &passCmd, &response);
if (DEBUG) printf("<PASS RESPONSE>\n%s\n", response);

if(strncmp("530", response, 3) == 0){
    printf("Login incorrect\n");
    exit(-1);
}
else if(strncmp("230", response, 3) != 0)
{
    printf("Login failed\n");
    exit(-1);
}

if (!DEBUG) printf("Login successful\n");

//--- PASV ---

if (DEBUG) printf("<PASV CMD>\nPASV\n\n");

received = sendComand(sd, "PASV\r\n", response);
if (DEBUG) printf("<PASV RESPONSE>\n%s\n", response);

if(strncmp(response, "227", 3) != 0){
    printf("PASV request not accepted\n");
    exit(-1);
}

//--- 2ND SOCKET ---

int i = 0;
char *token; //Pointer to hold each token parsed from the response.
int parsedIP[6];

```

```

        token = strtok(response, " (,)"); //parse the response by the
delimiters '(', ')' and ','

        while( (token = strtok(NULL, " (,)")) != NULL && i < 6)
{
    parsedIP[i] = atoi(token);
    //printf("Parsed value %d : %d\n", i, parsedIP[i]);
    i++;
}
if(i < 6){
    printf("Error parsing Pasv IP address.\n");
    exit(-1);
}

char pasvIpAddr[20];
sprintf(pasvIpAddr, "%d.%d.%d.%d", parsedIP[0], parsedIP[1],
parsedIP[2], parsedIP[3]);

struct sockaddr_in addr;
addr.sin_family = AF_INET;

addr.sin_port = htons( (parsedIP[4] << 8) + parsedIP[5] ); //cria a
porta para a socket multiplicando o byte 4 por 256 e somando o byte 5
inet_aton(pasvIpAddr, (struct in_addr*) &addr.sin_addr.s_addr);
//converte o IP para uma struct in_addr

if (DEBUG) printf("Parsed Pasv IP address - %s:%d\n\n", pasvIpAddr,
addr.sin_port);

int sd2 = socket(AF_INET, SOCK_STREAM, 0);
if( (connect(sd2, (const struct sockaddr*) &addr, sizeof(addr)) ) ==
-1){
    printf("Data connection not established.\n");
    exit(-1);
}

//--- TYPE ---

if (DEBUG) printf("<TYPE CMD>\nTYPE I\n\n");

received = sendComand(sd, "TYPE I\r\n", response);
if (DEBUG) printf("<TYPE RESPONSE>\n%s\n", response);

if(strncmp(response, "200", 3) != 0){
    printf("TYPE request not accepted\n");
    exit(-1);
}

```

```

//--- SIZE ---

char sizeCmd[512];
if(strlen(cInfo.filepath) > 0)
    sprintf(sizeCmd, "SIZE %s/%s\r\n", cInfo.filepath,
cInfo.filename);
else
    sprintf(sizeCmd, "SIZE %s\r\n", cInfo.filename);

if (DEBUG) printf("<SIZE CMD>\nSIZE\n\n");

received = sendComand(sd, &sizeCmd, &response);
if (DEBUG) printf("<SIZE RESPONSE>\n%s\n", response);

if(strncmp("550", response, 3) == 0){
    printf("No such file or directory.\n");
    exit(-1);
}
else if(strncmp("213", response, 3) != 0){
    printf("No such file or directory.\n");
    exit(-1);
}

int filesize = atoi(response + 4);

//--- RETR ---


char retrCmd[512];
if(strlen(cInfo.filepath) > 0)
    sprintf(retrCmd, "RETR %s/%s\r\n", cInfo.filepath,
cInfo.filename);
else
    sprintf(retrCmd, "RETR %s\r\n", cInfo.filename);

if (DEBUG) printf("<RETR CMD>\n%s\n", retrCmd);

received = sendComand(sd, &retrCmd, &response);
if (DEBUG) printf("<RETR RESPONSE>\n%s\n", response);

if(strncmp("150", response, 3) != 0 && strncmp("125", response, 3) != 0){
    printf("File is unavailable/not found.\n");
    exit(-1);
}

if (DEBUG) printf("Transfering file (bytes received)\n");
else printf("Transfering file...\n");

```

```

//--- Transfer file ---


dataConnection(sd2, cInfo.filename, filesize);

close(sd2);

//--- END ---


received = recv(sd, response, RESP_BUFFER_SIZE, 0);
response[received] = 0;
if (DEBUG) printf("\n<RETR FINAL RESPONSE>\n%s\n", response);

if(strncmp(response, "226", 3) != 0){
    printf("Transfer failed.\n");
    exit(-1);
}
else{
    printf("Transfer complete!\n");
}

close(sd);
}

void dataConnection(int sd2, char * filename, int filesize)
{
    char buffer[FILE_BUFFER_SIZE];
    int bytesReceived;
    int totalBytes = 0;

    int fd = open(filename, O_CREAT|O_TRUNC|O_WRONLY, 0777);

    while (1) {

        bytesReceived = recv(sd2, buffer, FILE_BUFFER_SIZE, 0);

        if (bytesReceived == 0) {
            break;
        } else if (bytesReceived == -1) {
            perror("Error receiving file data");
            break;
        } else {
            int bytesWritten = 0;
            while (bytesWritten < bytesReceived) {
                int result = write(fd, buffer + bytesWritten,
bytesReceived - bytesWritten);
                if (result == -1) {
                    perror("Error writing file data");
                    break;
                }
            }
        }
    }
}

```

```

        }
        bytesWritten += result;
    }
    totalBytes += bytesReceived;
}

if (DEBUG) {
    printf("(%d)", bytesReceived);
} else {
    loadingBar(totalBytes, filesize);
}

if (totalBytes == filesize) {
    break;
}
printf("\n");

close(fd);

if (totalBytes != filesize){
    printf("File size mismatch: %d (received) / %d (expected)\n",
totalBytes, filesize);
    //TODO delete file
    exit(1);
}
}

/*
ftp://user:pass@hostname:2121/path/to/file.txt, the function will extract:

user: "user"
password: "pass"
hostname: "hostname"
port: "2121"
filepath: "path/to"
filename: "file.txt"
*/
}

int parseUrl(char* url, connectionInfo *cInfo)
{
    int pos_arroba = -1;
    int first_bar = -1;
    int second_bar = -1;
    int first_colon = -1;
    int second_colon = -1;
    int third_bar = -1;
    int colon_count = 0;

```

```

int anonymous = 1;
unsigned int i;

// First cycle to determine if url has a valid syntax;
for (i = 0; i < strlen(url); i++) {
    if (url[i] == '@') {
        anonymous = 0;
        pos_arroba = i;
    }
    if (url[i] == '/') {
        if (first_bar == -1)
            first_bar = i;
        else if (second_bar == -1) {
            if (first_bar == i - 1)
                second_bar = i;
            else
                return -1;
        } else if (third_bar == -1)
            third_bar = i;
    }
    if (url[i] == ':' && first_bar != -1 && second_bar != -1) {
        if (colon_count > 1)
            return -3;
        else
            colon_count++;
    }
}

for (i = 0; i < strlen(url); i++) {
    if (url[i] == ':' && i > second_bar) {
        if (first_colon == -1)
            first_colon = i;
        else
            second_colon = i;
    }
}

int hostname_begin = 0;
if (anonymous == 0) {
    memcpy(cInfo->user, url + second_bar + 1, first_colon - second_bar -
1);
    cInfo->user[first_colon - second_bar - 1] = '\0';

    memcpy(cInfo->password, url + first_colon + 1, pos_arroba -
first_colon - 1);
    cInfo->password[pos_arroba - first_colon - 1] = '\0';

    hostname_begin = pos_arroba;
}

```

```

    } else {
        strcpy(cInfo->user, "ftp");
        strcpy(cInfo->password, "eu@");

        hostname_begin = second_bar;
    }

    if (second_colon != -1) {
        memcpy(cInfo->hostname, url + hostname_begin + 1, second_colon -
hostname_begin - 1);
        cInfo->hostname[second_colon - hostname_begin - 1] = '\0';

        memcpy(cInfo->port, url + second_colon + 1, third_bar - second_colon
- 1);
        cInfo->port[third_bar - second_colon - 1] = '\0';
    } else {
        memcpy(cInfo->hostname, url + hostname_begin + 1, third_bar -
hostname_begin - 1);
        cInfo->hostname[third_bar - hostname_begin - 1] = '\0';

        strcpy(cInfo->port, "21");
    }

    char *last_barPtr = strrchr(url, '/');
    int last_bar = last_barPtr - url;
    if (last_bar > third_bar) {
        memcpy(cInfo->filepath, url + third_bar + 1, last_bar - third_bar -
1);
        cInfo->filepath[last_bar - third_bar - 1] = '\0';
    } else {
        cInfo->filepath[0] = '\0';
    }

    memcpy(cInfo->filename, last_barPtr + 1, strlen(url) - last_bar - 1);
    cInfo->filename[strlen(url) - last_bar - 1] = '\0';

    return 0;
}

void loadingBar(float file_size_processed, float file_total_size)
{
    float percentage = 100.0 * file_size_processed / file_total_size;
    printf("\rStatus: %6.2f% [", percentage);      //display progress from
0 to 100%

    int i, len = 50;
    int pos = percentage * len / 100.0;

```

```

//displays visual progress
for (i = 0; i < len; i++){
    if(i <= pos)
        printf("=");
    else
        printf(" ");
}

printf("]");

fflush(stdout);
}

```

Configuration Commands:

Experiment 1 Commands

tux13:
ifconfig eth1 172.16.10.1/24
tux14:
ifconfig eth1 172.16.10.254/24
tux13:
ping 172.16.10.254
tux14:
ping 172.16.10.1
tux13:
arp -a
arp -d 172.16.10.254/24
arp -a

Experiment 2 Commands

Switch console:
/system reset-configuration
tux12:
ifconfig eth1 172.16.11.1/24
Switch console:
/interface bridge add name=bridge50
/interface bridge add name=bridge51
/interface bridge port remove [find interface=ether4]
/interface bridge port remove [find interface=ether7]
/interface bridge port remove [find interface=ether12]
/interface bridge port add bridge=bridge0 interface=ether4
/interface bridge port add bridge=bridge0 interface=ether7
/interface bridge port add bridge=bridge1 interface=ether12

```
tux13:  
    ping 172.16.10.254  
    ping 172.16.11.1  
    ping -b 172.16.10.255  
tux12:  
    ping -b 172.16.11.255
```

Experiment 3 Commands

```
tux14:  
    ifconfig eth2 up  
    ifconfig eth2 172.16.11.253/24
```

Switch console:

```
/interface bridge port remove [find interface=ether15]  
/interface bridge port add bridge=bridge1 interface=ether15
```

```
tux14:  
    sysctl net.ipv4.ip_forward=1  
    sysctl net.ipv4.icmp_echo_ignore_broadcasts=0
```

```
tux12:  
    route add -net 172.16.10.0/24 gw 172.16.11.253
```

```
tux13:  
    route add -net 172.16.11.0/24 gw 172.16.10.254  
    ping 172.16.10.254  
    ping 172.16.11.253  
    ping 172.16.11.1
```

```
tux12:  
    arp -d 172.16.11.253
```

```
tux13:  
    arp -d 172.16.10.254
```

```
tux14:  
    arp -d 172.16.10.1  
    arp -d 172.16.11.1
```

Experiment 4 Commands

Switch console:

```
/interface bridge port remove [find interface=ether24]  
/interface bridge port add bridge=bridge1 interface=ether24
```

Router console: /system reset-configuration

```
/ip address add address=172.16.1.11/24 interface=ether1  
/ip address add address=172.16.11.254/24 interface=ether2  
/ip route add dst-address=172.16.10.0/24 gateway=172.16.11.253  
/ip route add address=0.0.0.0/0 gateway=172.16.1.254
```

```
tux12:
```

```

        route add default gw 172.16.11.254
tux13:
        route add default gw 172.16.10.254
tux14:
        route add default gw 172.16.11.254
tux13:
        ping 172.16.10.254
        ping 172.16.11.1
        ping 172.16.11.254
tux12:
        sysctl net.ipv4.conf.eth1.accept_redirects=0
        sysctl net.ipv4.conf.all.accept_redirects=0
        route del -net 172.16.10.0 gw 172.16.11.253/24
        ping 172.16.10.1
        traceroute -n 172.16.10.1
        route add -net 172.16.10.0/24 gw 172.16.11.253
        traceroute -n 172.16.10.1
        sysctl net.ipv4.conf.eth1.accept_redirects=1
        sysctl net.ipv4.conf.all.accept_redirects=1
tux13:
        ping 172.16.1.10
Router console:
        /ip firewall nat disable 0
tux13:
        ping 172.16.1.10
Router console:
        /ip firewall nat enable 0

```

Experiment 5 Commands

add the DNS server (10.227.20.3) in case it isn't seated as default.
Tux12, Tux13, Tux14: ping google.com

Experiment 6 Commands

run the ftp application in the tux13

Wireshark Captures:

Experiment 1 Captures

1.1 - Ping from tux13 to tux14

5 2. 500886399	EdimaxTe_ed:bb:37	Broadcast	ARP	42 Who has 172.16.10.254? Tell 172.16.10.1
6 2. 500907738	Kali-04:20:99	EdimaxTe_ed:bb:37	ARP	60 172.16.10.254. at 08:41:0f:44:2b:99
7 2. 500969490	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x1644, seq=1/256, ttl=64 (reply in 8)
8 2. 510062728	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x1644, seq=1/256, ttl=64 (request in 7)
9 3. 527947975	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x1644, seq=2/512, ttl=64 (reply in 10)
10 3. 527151487	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x1644, seq=2/512, ttl=64 (request in 9)
11 4. 604374359	Routerbo_1c:8c:9f	Spanning-tree-(for...	STP	60 RST. Root = 32768/0/c4:ad:34:1c:8c:99 Cost = 0 Port = 0x8007
12 4. 604374348	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x1644, seq=3/768, ttl=64 (reply in 13)
13 4. 551131392	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x1644, seq=3/768, ttl=64 (request in 12)
14 5. 575051582	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x1644, seq=4/1024, ttl=64 (reply in 15)
15 5. 575138464	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x1644, seq=4/1024, ttl=64 (request in 14)

Experiment 2 Captures

2.1 - Ping from tux13 to tux14

```

12 14.717864923 172.16.10.1 172.16.10.254 ICMP 98 Echo (ping) request id=0x6490, seq=1/256, ttl=64 (reply in 13)
13 14.719018294 172.16.10.254 172.16.10.1 ICMP 98 Echo (ping) reply id=0x6490, seq=1/256, ttl=64 (request in 12)
14 15.737047500 172.16.10.1 172.16.10.254 ICMP 98 Echo (ping) request id=0x6490, seq=2/512, ttl=64 (reply in 15)
15 15.737182642 172.16.10.254 172.16.10.1 ICMP 98 Echo (ping) reply id=0x6490, seq=2/512, ttl=64 (request in 14)

```

2.2 - Ping from tux13 to tux12

```

21 34.65211543 Routerboard_1c:8c:a4 Spanning-tree-(for-br- 80 RST, Root = 32768/0/c4:ad:34:1c:8c:a4 Cost = 0 Port = 0x8001
22 34.392052968 172.16.11.1 172.16.11.255 ICMP 98 Echo (ping) request id=0x7ea2, seq=4/1024, ttl=64 (no response found!)
23 35.416938537 172.16.11.1 172.16.11.255 ICMP 98 Echo (ping) request id=0x7ea2, seq=5/1280, ttl=64 (no response found!)
24 36.040032638 Routerboard_1c:8c:a4 Spanning-tree-(for- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:a4 Cost = 0 Port = 0x8001
25 36.440058328 172.16.11.1 172.16.11.255 ICMP 98 Echo (ping) request id=0x7ea2, seq=6/1536, ttl=64 (no response found!)
26 37.464834259 172.16.11.1 172.16.11.255 ICMP 98 Echo (ping) request id=0x7ea2, seq=7/1792, ttl=64 (no response found!)
27 38.488037637 Routerboard_1c:8c:a4 Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:a4 Cost = 0 Port = 0x8001
28 38.488037637 172.16.11.1 172.16.11.255 ICMP 98 Echo (ping) request id=0x7ea2, seq=8/2048, ttl=64 (no response found!)
29 39.512935708 172.16.11.1 172.16.11.255 ICMP 98 Echo (ping) request id=0x7ea2, seq=9/2304, ttl=64 (no response found!)

```

Experiment 3 Captures

3.1.1 - ICMP packets, captured in tux13 during the ping from tux13 to the eth1 interface of tux14

```

11 0.027662 172.16.10.1 172.16.10.254 ICMP 98 Echo (ping) request id=0xeeee, seq=2/512, ttl=64 (reply in 12)
12 0.027779 172.16.10.254 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=2/512, ttl=64 (request in 11)
13 0.028065 172.16.10.1 172.16.10.254 ICMP 98 Echo (ping) request id=0xeeee, seq=3/768, ttl=64 (reply in 14)
14 0.028036 172.16.10.254 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=3/768, ttl=64 (request in 13)
15 0.028152 Routerboard_1c:8c:9f Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:9f Cost = 0 Port = 0x8002
16 0.028390 172.16.10.1 172.16.10.254 ICMP 98 Echo (ping) request id=0xeeee, seq=4/1024, ttl=64 (reply in 17)
17 0.028489 172.16.10.254 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=4/1024, ttl=64 (request in 16)

```

3.1.2 - ICMP packets, captured in tux13 during the ping from tux13 to the eth2 interface of tux14

```

24 0.029242 172.16.10.1 172.16.11.253 ICMP 98 Echo (ping) request id=0xeeee, seq=1/256, ttl=64 (reply in 25)
25 0.029403 172.16.11.253 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=1/256, ttl=64 (request in 24)
26 0.029583 172.16.10.1 172.16.11.253 ICMP 98 Echo (ping) request id=0xeeee, seq=2/512, ttl=64 (reply in 27)
27 0.029796 Routerboard_1c:8c:9f Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:9f Cost = 0 Port = 0x8002
28 0.030054 172.16.10.1 172.16.11.253 ICMP 98 Echo (ping) request id=0xeeee, seq=3/768, ttl=64 (reply in 30)
29 0.030235 172.16.11.253 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=3/768, ttl=64 (request in 29)
30 0.030368 172.16.10.1 172.16.11.253 ICMP 98 Echo (ping) request id=0xeeee, seq=4/1024, ttl=64 (reply in 32)
31 0.030538 172.16.11.253 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=4/1024, ttl=64 (request in 31)

```

3.1.3 - ICMP packets, captured in tux13 during the ping from tux13 to tux12

```

35 0.030683 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0xeeee, seq=1/256, ttl=64 (reply in 36)
36 0.031020 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=1/256, ttl=64 (request in 35)
37 0.031155 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0xeeee, seq=2/512, ttl=64 (reply in 38)
38 0.031276 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=2/512, ttl=64 (request in 37)
39 0.031480 Routerboard_1c:8c:9f Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:9f Cost = 0 Port = 0x8002
40 0.031515 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0xeeee, seq=3/768, ttl=64 (reply in 41)
41 0.031666 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=3/768, ttl=64 (request in 40)
42 0.031807 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0xeeee, seq=4/1024, ttl=64 (reply in 43)
43 0.031951 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0xeeee, seq=4/1024, ttl=64 (request in 42)

```

3.2.1 - ARP and ICMP packets, captured in tux14's eth1 during the establishment of the connection between tux13 and tux12

```

21 25.007471767 EdimaxTechno_ed:bb:37 Broadcast ARP 60 Who has 172.16.10.254? Tell 172.16.10.1
22 25.007474620 KVM_04:20:99 EdimaxTechno_ed:bb:37 ARP 42 172.16.10.254 is at 00:00:00:00:00:00
23 25.067601744 172.16.11.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=1/256, ttl=64 (reply in 24)
24 25.067885586 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) request id=0x068b, seq=1/256, ttl=64 (request in 23)
25 26.029454252 Routerboard_1c:8c:9c Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:9c Cost = 0 Port = 0x8001
26 26.030079509 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=2/512, ttl=64 (reply in 27)
27 26.030292095 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0x068b, seq=2/512, ttl=64 (request in 26)
28 27.116284139 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=3/768, ttl=64 (reply in 29)
29 27.116435907 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0x068b, seq=3/768, ttl=64 (request in 28)
30 28.011708652 Routerboard_1c:8c:9c Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:9c Cost = 0 Port = 0x8001
31 28.140204356 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=4/1024, ttl=64 (reply in 32)

```

3.2.2 - ARP e ICMP packets, captured in tux14's eth2 during the establishment of the connection between tux13 and tux12

```

18 23.065282338 TpLinkTechno_02:03:02 Broadcast ARP 42 Who has 172.16.11.17? Tell 172.16.11.253
19 23.065400722 Netronix_5b:35:0c TpLinkTechno_02:03:02 ARP 60 172.16.11.1 is at 00:00:54:59:35:0c
20 23.065417973 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=1/256, ttl=64 (reply in 21)
21 23.065540617 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0x068b, seq=1/256, ttl=64 (request in 20)
22 23.065540617 Routerboard_1c:8c:a7 Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:a7 Cost = 0 Port = 0x8001
23 24.089937155 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=2/512, ttl=64 (reply in 24)
24 24.090044783 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0x068b, seq=2/512, ttl=64 (request in 23)
25 25.113973393 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=3/768, ttl=64 (reply in 26)
26 25.114083116 172.16.11.1 172.16.10.1 ICMP 98 Echo (ping) reply id=0x068b, seq=3/768, ttl=64 (request in 25)
27 25.114083116 Routerboard_1c:8c:a7 Spanning-tree-(for-br- 60 RST, Root = 32768/0/c4:ad:34:1c:8c:a7 Cost = 0 Port = 0x8001
28 26.137969070 172.16.10.1 172.16.11.1 ICMP 98 Echo (ping) request id=0x068b, seq=4/1024, ttl=64 (reply in 29)

```

Experiment 4 Captures

4.1.1 - ICMP packets, captured in tux13 during the ping from tux13 to the eth1 interface of tux14

11 9. 052344519	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x0675, seq=1/256, ttl=64 (reply in 12)
12 9. 052486790	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0675, seq=1/256, ttl=64 (request in 11)
13 10. 016180544	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 RST Root = 32768/0/c4:ad:34:1c:8c:9c Cost = 0 Port = 0x0002
14 10. 058756186	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x0675, seq=2/512, ttl=64 (reply in 15)
15 10. 058756185	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0675, seq=2/512, ttl=64 (request in 14)
16 11. 092760428	172.16.10.1	172.16.10.254	ICMP	98 Echo (ping) request id=0x0675, seq=3/768, ttl=64 (reply in 17)
17 11. 092876573	172.16.10.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0675, seq=3/768, ttl=64 (request in 16)

4.1.2 - ICMP packets, captured in tux13 during the ping from tux13 to tux12

46 46. 1640904243	172.16.10.1	172.16.11.1	ICMP	98 Echo (ping) request id=0x068b, seq=1/256, ttl=64 (reply in 47)
47 46. 164296669	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) reply id=0x068b, seq=1/256, ttl=63 (request in 46)
48 47. 178749826	172.16.10.1	172.16.11.1	ICMP	98 Echo (ping) request id=0x068b, seq=2/512, ttl=64 (reply in 49)
49 48. 050558724	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) reply id=0x068b, seq=2/512, ttl=63 (request in 48)
50 48. 050558691	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 RST Root = 32768/0/c4:ad:41:1c:8c:9c Cost = 0 Port = 0x0002
51 48. 202752384	172.16.10.1	172.16.11.1	ICMP	98 Echo (ping) request id=0x068b, seq=3/768, ttl=64 (reply in 52)
52 49. 226766381	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) reply id=0x068b, seq=3/768, ttl=63 (request in 51)
53 49. 226766384	172.16.10.1	172.16.11.1	ICMP	98 Echo (ping) request id=0x068b, seq=4/1024, ttl=64 (reply in 54)
54 49. 227001470	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) reply id=0x068b, seq=4/1024, ttl=63 (request in 53)

4.1.3 - ICMP packets, captured in tux13 during the ping from tux13 to the router

66 59. 500253796	172.16.10.1	172.16.11.254	ICMP	98 Echo (ping) request id=0x0696, seq=1/256, ttl=64 (reply in 67)
67 59. 500558724	172.16.11.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0696, seq=1/256, ttl=63 (request in 66)
68 60. 522755181	172.16.10.1	172.16.11.254	ICMP	98 Echo (ping) request id=0x0696, seq=2/512, ttl=64 (reply in 70)
70 60. 523061581	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0696, seq=2/512, ttl=63 (request in 69)
73 61. 549077	172.16.10.1	172.16.11.254	ICMP	98 Echo (ping) request id=0x0696, seq=3/768, ttl=64 (reply in 72)
75 62. 047000981	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0696, seq=3/768, ttl=63 (request in 71)
73 62. 050000083	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 RST Root = 32768/0/c4:ad:34:1c:8c:9c Cost = 0 Port = 0x0002
74 62. 570758929	172.16.10.1	172.16.11.254	ICMP	98 Echo (ping) request id=0x0696, seq=4/1024, ttl=64 (reply in 75)
75 63. 571321122	172.16.11.254	172.16.10.1	ICMP	98 Echo (ping) reply id=0x0696, seq=4/1024, ttl=63 (request in 74)
76 63. 594758446	172.16.10.1	172.16.11.254	ICMP	98 Echo (ping) request id=0x0696, seq=5/1280, ttl=64 (request in 77)

4.2 - ICMP packets, captured in tux12 during the ping from tux12 to tux13 after removing the route between tux12 and tux14

36 44. 276322692	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) request id=0x07bf, seq=3/768, ttl=64 (reply in 38)
37 44. 276470896	172.16.11.254	172.16.11.1	ICMP	126 Redirect (Redirect for host)
38 44. 276656117	172.16.10.1	172.16.11.1	ICMP	98 Echo (ping) reply id=0x07bf, seq=3/768, ttl=63 (request in 36)
39 45. 300322018	172.16.11.1	172.16.10.1	ICMP	98 Echo (ping) request id=0x07bf, seq=4/1024, ttl=64 (reply in 41)
40 45. 300481948	172.16.11.254	172.16.11.1	ICMP	126 Redirect (Redirect for host)
41 45. 300663614	172.16.10.1	172.16.11.1	ICMP	98 Echo (ping) reply id=0x07bf, seq=4/1024, ttl=63 (request in 39)

4.3 - ICMP packets, captured in tux13 during the ping from tux13 to the server FTP(172.16.1.10) with NAT enabled

8 4. 949735664	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) request id=0x109, seq=3/768, ttl=64 (reply in 0)
9 4. 950137256	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) reply id=0x109, seq=3/768, ttl=62 (request in 8)
10 5. 973736117	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) request id=0x109, seq=4/1024, ttl=64 (reply in 11)
11 5. 974143566	172.16.1.10	172.16.1.10	ICMP	98 Echo (ping) reply id=0x109, seq=4/1024, ttl=62 (request in 10)
12 6. 999735676	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 RST Root = 32768/0/c4:ad:34:1c:8c:9c Cost = 0 Port = 0x0002
13 6. 999735676	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) request id=0x109, seq=5/1280, ttl=64 (reply in 14)
14 6. 999824445	172.16.1.10	172.16.1.10	ICMP	98 Echo (ping) reply id=0x109, seq=5/1280, ttl=62 (request in 13)
15 7. 021730702	172.16.1.10	172.16.1.10	ICMP	98 Echo (ping) request id=0x109, seq=6/1536, ttl=64 (reply in 17)
16 8. 0221245832	172.16.1.10	172.16.1.10	ICMP	98 Echo (ping) reply id=0x109, seq=6/1536, ttl=62 (request in 16)
17 8. 045730501	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) request id=0x109, seq=7/1792, ttl=64 (reply in 19)
19 9. 046091646	172.16.1.10	172.16.10.1	ICMP	98 Echo (ping) reply id=0x109, seq=7/1792, ttl=62 (request in 18)

4.4 - ICMP packets, captured in tux13 during the ping from tux13 to the server FTP(172.16.1.10) with NAT disabled

11 8. 066859994	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) request id=0x1197, seq=7/1792, ttl=64 (no response found!)
12 8. 092862758	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 RST Root = 32768/0/c4:ad:34:1c:8c:9c Cost = 0 Port = 0x0002
14 11. 016861695	172.16.10.1	172.16.1.10	ICMP	98 Echo (ping) request id=0x1197, seq=9/2304, ttl=64 (no response found!)
15 12. 040866932	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 Echo (ping) request id=0x1197, seq=10/2560, ttl=64 (no response found!)
17 14. 084572493	Routerboard0_ic:8c:9f	Spanning-tree-(for-brdg_	STP	98 Echo (ping) request id=0x1197, seq=11/2816, ttl=64 (no response found!)

Experiment 5 Captures

5.1 - DNS and ICMP packets, captured in tux13 during the ping from tux13 to google.com

46 4. 373766885	10.227.20.13	10.227.20.3	DNS	70 Standard query 0xcfcfa A google.com
47 4. 373719772	10.227.20.13	10.227.20.3	DNS	70 Standard query 0xdaf24 AAAA google.com
48 4. 374278988	10.227.20.3	10.227.20.13	DNS	88 Standard query response 0xcfcfa google.com A 142.250.185.14
49 4. 374295191	10.227.20.3	10.227.20.13	DNS	98 Standard query response 0xdaf24 AAAA google.com AAAA 2a0:1450:4003:80f::200e
50 4. 374589151	10.227.20.13	142.250.185.14	ICMP	98 Echo (ping) request id=0x13e9, seq=1/256, ttl=64 (reply in 51)
51 4. 392523383	142.250.185.14	10.227.20.13	ICMP	98 Echo (ping) reply id=0x13e9, seq=1/256, ttl=112 (request in 50)

Experiment 6 Captures

6.1 - TCP and FTP packets, captured in tux13 during the download of a file in the FTP server

8 7.394604768	172.16.10.1	172.16.1.10	TCP	74 38956 21 [SYN] Seq=0 Win=64249 Len=0 MSS=1460 SACK_PERM TSval=2588564162 TSecr=157538789
7.3951059022	172.16.1.10	172.16.10.1	TCP	74 21 → 38956 [SYN ACK] Seq=0 Ack=1 Win=65168 Len=0 MSS=1460 SACK_PERM TSval=157538549
7.395148875	172.16.10.1	172.16.1.10	TCP	116 Response: 220 ProFTPD Server (Debian) [::ffff:172.16.1.10]
9.6269933360	172.16.1.10	172.16.10.1	TCP	66 38956 → 21 [ACK] Seq=1 Ack=21 Win=64256 Len=0 TSval=2588564183 TSecr=157538549
10.627021111	172.16.1.10	172.16.10.1	TCP	66 38956 → 21 [ACK] Seq=1 Ack=21 Win=64256 Len=0 TSval=2588564183 TSecr=157538789
10.627021252	172.16.1.10	172.16.1.10	FTP	77 Request: USER rcom
11.627281410	172.16.1.10	172.16.1.10	TCP	66 21 → 38956 [ACK] Seq=51 Ack=12 Win=64256 Len=0 TSval=157538781 TSecr=258856444...
12.627281429	172.16.1.10	172.16.1.10	TCP	99 Response: 331 User name required for rcom
13.669729296	172.16.1.10	172.16.1.10	TCP	66 38956 → 21 [ACK] Seq=12 Ack=83 Win=64256 Len=0 TSval=2588564457 TSecr=1575387...
Rerouting traffic (for bridge)				
15.8.627188963	172.16.1.10	172.16.1.10	FTP	77 Request: PASS rcom
16.8.668972041	172.16.1.10	172.16.1.10	TCP	66 21 → 38956 [ACK] Seq=83 Ack=23 Win=65280 Len=0 TSval=157539823 TSecr=25885654...
17.8.805457861	172.16.1.10	172.16.1.10	FTP	112 Response: 230-Welcome, archive user rcom@172.16.1.10
18.8.805482693	172.16.1.10	172.16.1.10	FTP	66 38956 → 21 [ACK] Seq=83 Ack=12 Win=64256 Len=0 TSval=2588565593 TSecr=157539...
19.8.805482515	172.16.1.10	172.16.1.10	FTP	69 Response:
20.8.805487543	172.16.1.10	172.16.1.10	FTP	66 38956 → 21 [ACK] Seq=23 Ack=132 Win=64256 Len=0 TSval=2588565593 TSecr=157539...
21.8.805487523	172.16.1.10	172.16.1.10	FTP	112 Response: 230-The local time is: Thu Dec 19 17:26:49 2024
22.8.805514153	172.16.1.10	172.16.1.10	TCP	66 38956 → 21 [ACK] Seq=23 Ack=178 Win=64256 Len=0 TSval=2588565593 TSecr=157539...
23.8.805514070	172.16.1.10	172.16.1.10	FTP	142 Response:
24.8.805514720	172.16.1.10	172.16.1.10	TCP	66 38956 → 21 [ACK] Seq=23 Ack=254 Win=64256 Len=0 TSval=2588565593 TSecr=157539...
25.8.805615632	172.16.1.10	172.16.1.10	FTP	132 Response: please report them via e-mail to <root@ftp.netlab.fe.up.pt>
26.8.805620381	172.16.1.10	172.16.1.10	FTP	66 38956 → 21 [ACK] Seq=232 Win=64256 Len=0 TSval=2588565593 TSecr=157539...
27.8.805620360	172.16.1.10	172.16.1.10	FTP	91 Response: 230 User rcom login in
28.8.805770540	172.16.1.10	172.16.1.10	FTP	66 38956 → 21 [ACK] Seq=23 Ack=345 Win=64256 Len=0 TSval=2588565593 TSecr=157539...
29.9.627349329	172.16.1.10	172.16.1.10	FTP	72 Request: PASV
30.9.627349342	172.16.1.10	172.16.1.10	FTP	66 38956 → 21 [ACK] Seq=345 Ack=29 Win=65280 Len=0 TSval=157540781 TSecr=258856...
31.9.628362244	172.16.1.10	172.16.1.10	FTP	116 Response: 227 Entering Passive Mode (172.16.1.10,159,125)
32.9.628368739	172.16.1.10	172.16.1.10	TCP	66 38956 → 21 [ACK] Seq=299 Ack=395 Win=64256 Len=0 TSval=2588566416 TSecr=157540...
Rerouting traffic (for bridge)				
34.10.6272520815	172.16.10.1	172.16.1.10	TCP	74 47098 → 40829 [SYN] Seq=0 Win=64249 Len=0 MSS=1460 SACK_PERM TSval=2588567415...
35.10.6272520815	172.16.10.1	172.16.1.10	TCP	74 47098 → 40829 [SYN ACK] Seq=1 Ack=1 Win=64256 Len=0 MSS=1460 SACK_PERM TSval=157541...
36.10.628899214	172.16.1.10	172.16.1.10	TCP	66 47098 → 40829 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TSval=2588567416 TSecr=157541...
37.10.628899531	172.16.10.1	172.16.1.10	FTP	129 Response: 150 Opening BINARY mode data connection for README (58 bytes)
38.10.628733377	172.16.1.10	172.16.1.10	FTP	66 47098 → 40829 [ACK] Seq=232 Win=64256 Len=0 TSval=2588569417 TSecr=157543...
39.10.628733197	172.16.1.10	172.16.1.10	FTP	74 Request: TYPE I
40.11.628182585	172.16.1.10	172.16.1.10	FTP	66 47098 → 40829 [ACK] Seq=232 Win=64256 Len=0 TSval=2588567416 TSecr=157541...
41.11.629060654	172.16.1.10	172.16.1.10	FTP	79 Request: SIZE README
42.11.629918154	172.16.1.10	172.16.1.10	TCP	74 Response: 213 58
Rerouting traffic (for bridge)				
44.12.589892332	KYE 04:29:09	EdimaxTechno_ed:bb:37	ARP	60 Who has 172.16.10.1? Tell 172.16.10.254
45.12.6272520815	172.16.10.1	172.16.10.1	ARP	42 172.16.10.1 is the source port 50:fc:c6:bb:37
46.12.628322338	172.16.10.1	172.16.1.10	FTP	79 Request: RETR README
47.12.629500607	172.16.1.10	172.16.1.10	FTP	129 Response: 150 Opening BINARY mode data connection for README (58 bytes)
48.12.629500609	172.16.1.10	172.16.1.10	FTP	66 47098 → 40829 [ACK] Seq=232 Win=64256 Len=0 TSval=2588569417 TSecr=157543...
49.12.629615726	172.16.1.10	172.16.1.10	FTP-DATA	124 FTP Data: 58 bytes (TYPE I)
50.12.629622719	172.16.1.10	172.16.1.10	TCP	66 47098 → 40829 [ACK] Seq=1 Ack=59 Win=64256 Len=0 TSval=2588569417 TSecr=15754...
51.12.629622719	172.16.1.10	172.16.1.10	FTP	66 47098 → 40829 [ACK] Seq=1 Ack=59 Win=64256 Len=0 TSval=2588569417 TSecr=15754...
52.12.669724432	172.16.10.1	172.16.1.10	TCP	66 47098 → 40829 [ACK] Seq=1 Ack=60 Win=64256 Len=0 TSval=2588569457 TSecr=15754...
53.13.628602669	172.16.1.10	172.16.1.10	TCP	66 47098 → 40829 [ACK] Seq=1 Ack=60 Win=64256 Len=0 TSval=2588570418 TSecr=15754...
54.13.628602671	172.16.1.10	172.16.1.10	FTP	66 47098 → 40829 [ACK] Seq=1 Ack=60 Win=65280 Len=0 TSval=157544783 TSecr=258857...
55.13.629436801	172.16.1.10	172.16.1.10	FTP	89 Response: 226 Transfer complete
56.13.629445741	172.16.1.10	172.16.1.10	TCP	66 38956 → 21 [ACK] Seq=608 Win=64256 Len=0 TSval=2588570418 TSecr=157544...
57.13.629910546	172.16.1.10	172.16.1.10	FTP	66 47098 → 40829 [ACK] Seq=1 Ack=608 Win=64256 Len=0 TSval=157544784 TSecr=25...
58.13.629910953	172.16.1.10	172.16.1.10	FTP	66 21 → 38956 [FIN ACK] Seq=608 Win=64256 Len=0 TSval=157544784 TSecr=25...
59.13.629924712	172.16.1.10	172.16.1.10	TCP	66 38956 → 21 [ACK] Seq=644 Ack=569 Win=64256 Len=0 TSval=2588570418 TSecr=15754...

Extra Annexes:

Download Application with the debug mode on

```
• pchmelo@pchmelo-VirtualBox:~/Downloads/RCom-project2-master$ ./download ftp://demo:password@test.rebex.net/readme.txt
<PARSE URL>
user = demo
password = password
hostname = test.rebex.net
port = 21
filepath =
filename = readme.txt

<CONNECT RESPONSE>
220-Welcome to test.rebex.net!
See https://test.rebex.net/ for more information and terms of use.
220 If you don't have an account, log in as 'anonymous' or 'ftp'.

<USER CMD>
USER demo

<USER RESPONSE>
331 Anonymous login OK, send your complete email address as your password.

<PASS CMD>
PASS password

<PASS RESPONSE>
230 User 'demo' logged in.

<PASV CMD>
PASV

<PASV RESPONSE>
227 Entering Passive Mode (194,108,117,16,4,26)

Parsed Pasv IP address - 194.108.117.16:5124

<TYPE CMD>
TYPE I

<TYPE RESPONSE>
200 TYPE set to I.

<SIZE CMD>
SIZE

<SIZE RESPONSE>
213 379

<RETR CMD>
RETR readme.txt

<RETR RESPONSE>
125 Data connection already open; starting 'BINARY' transfer.

Transferring file (bytes received)
(379)

<RETR FINAL RESPONSE>
226 Transfer complete.

Transfer complete!
```

Graph - Speed of packets forwarded by the network per second, in a single transfer of Tux53

